
COS 350: Program #6: multiple file line counting

Objectives: Working with fork, exec, pipes and comparing with pthreads and shared memory communication.

You will write 3 versions of a multiple file line counting program. Your goal is to make all three versions as fast as possible. Do not use `getc()` as in the class examples. Use `read()` to read in a large chunk of the file into a buffer for processing. This will be faster.

Version 1: lc1

Works like `wc` (word count) but simply counts the number of lines in each file specified on the command line, and for each file it prints the count and the filename, and finally the total. (Just count the number of newline characters. Your results should match `wc`.)

Version 2: lc2

Use `fork()` and `exec()` to create a separate process for each individual file. Just `exec()` `lc1` to do the counting and setup pipes to read back the output so that the main program can total them up.

The output should be identical to `lc1` except the output order may vary. Fork them all off so that they are working simultaneously.

Version 3: lc3

Use `pthread`s to count the files. Create simultaneous threads, one for each file. You can decide what sort of synchronization you need. As with `lc2`, the output order may vary. Remember to use the `-pthread` compiler flag.

Testing:

Once you have your programs all working, recompile them using the optimization flag `-O2` (e.g. `gcc -O2 ...`). This will make them run faster.

Use the `unix` program time for timing your programs. Time each of these running on the set of 10 files: `f1` – `f10` (these are in the course directory). Do the timing tests on an unloaded machine. You can check the load on a machine with the program: [uptime](#). The load values should be < 1.0 . Preferably close to 0.0.

Run at least 3 trials for each version and report the best time among these trials. Which version is fastest? Try to guess why.

What to turn in:

Written report:

1. Your name(s) and if a team, who did the electronic submit.
2. A discussion of any incomplete parts, known bugs, deviations from the specification, or extra features in your program.
3. Your code
4. Your Makefile
5. A log of the following testing results. (I recommend using script to capture it.)

```
time lc1 f*           test lc1 counting all 10 files f1-f10
time lc2 f*           test lc2
time lc3 f*           test lc3
exit                  exit script
```

6. A neatly organized table of your timing results like this:

Machine that the timing tests were performed on: _____

Try to use an unloaded machine and report the best of at least 3 trials for each program.

Timings results	real	user	sys
lc1			
lc2			
lc3			

7. Discuss which version was the fastest, and why you think that was?

Electronic submission:

- Before you submit your program, compile and test it on a linux machine in the lab.
- From a machine in the lab run the program “submit” to submit your files.
- Submit your source code (lc1.c lc2.c lc3.c) and any other source files, a Makefile that works for all 3 files, and your executables (lc1 lc2 lc3) to a directory named: prog6
- Do not in any way combine, compress, zip, or tar your files!

Grading:

- 5 points: Makefile
- 5 points: lc1 gets correct individual file line counts
- 5 points: lc1 gets correct total
- 15 points: lc2 the processes execute concurrently, and it is faster than lc1
- 15 points: lc2 gets correct total
- 15 points: lc3 the threads execute concurrently, and it is faster than lc1
- 15 points: lc3 gets correct total
- 10 points: Timing results and discussion
- 15 points: all versions finish in under 1 second